

(Copy from [here](#))

Custom Classifier Design Doc

Authors: zacharias.fisches@, vasily.vitchevsky@

[Description](#)

[Overview](#)

[Frontend](#)

[Backend](#)

[Details](#)

[Frontend](#)

[Navigation](#)

[Front page](#)

[TODO](#)

[Model editor](#)

[TODO](#)

[Backend](#)

[Web server](#)

[Models](#)

[Jobs / Tasks](#)

[train_model](#)

[run_model](#)

[Saving model state](#)

[Predefined Models](#)

[Linear Regression](#)

[Usage](#)

[Other](#)

[Redis](#)

[Possible use case](#)

Description

This document describes a service where a user can train and run a classifier on their data. By *user* we mean a data analyst who can write SQL queries, but doesn't have access to ML tools such as linear regression.

Eventually, a user will provide input data (e.g. SQL query result), choose a classifier (e.g. linear regression), train it, and finally use the classifier on yet another input data (e.g. SQL query result). The result will be saved to a new table.

As the input data is updated roughly every week, we have to presumably enable automatic weekly re-runs of the training of the model.

Overview

Frontend

The website is where the user works and submits train/predict requests.

The website contains following pages:

1. Page 1: List of all models, their status. Have train / predict buttons. Later - stop button.
2. Page 2: Create / Modify a model.
 - a. Model Name
 - b. Description
 - c. Train SQL query for training data
 - d. Dropdown: classifier selection
 - e. Eval SQL query for evaluation data

Backend

Train/predict requests from the frontend are sent to a backend (right now: the web server). The backend sends long running jobs to workers.

The result of a *training* job is a trained class instance. We save (pickle) the result to a models DB, along with the model parameters.

The result of a *predict* job is the predictions. We save them to a new table in the DB.

Details

Frontend

One solution would be to integrate with *redash*, but we then may have to modify our code when a new redash version is released, and it's probably more difficult. Therefore, for the first version we have a custom website.

Navigation

A top bar. Both "Soft-Query MVP" and "Home" brings us the the front page. "Create New Model" brings us to a new model page, which is similar to a "Modify model" page.

Front page

The front page contains the list of modules. Each model we can train, predict and modify. We also see the result of the last train/predict run:

Existing Soft-Query Models

Model:	Train	Run Model
test	Model	

Results:	stdout:	stderr:
<input type="text" value=":("/>	<input type="text"/>	<pre>ERROR:root:inconsistent use of tabs and spaces in indentation Traceback (most recent call last): File "./app/tasks.py", line 50, in train_model exec(model.ml_model, names) File "<string>", line 6 self.m = LinearRegression() ^ TabError: inconsistent use of tabs and spaces in indentation</pre>

TODO

- “Stop” button
- “Delete Model” button (perhaps better through the model edit page)
- “Duplicate Model” button (perhaps better through the model edit page)
- Style the page nicely
- Collapse long outputs so it doesn’t clutter the page.

Model editor

The model editor allows modifying a model:

- Name
- Description (multiline)
- Model code (python), with an option to load an existing model
- SQL query for train
- SQL query for prediction
- Save button

Model Editor

Model Name

test

Short Description

Choose an existing Python template to start (this will delete your current code!)

No Template: Loading this template will have no effect.

Load chosen model in the editor below

Model Python Code

```
1 import numpy as np
2 from sklearn.linear model import LinearRegression
```

SQL query to generate the train data

```
1 select price, quality from houses where quality is not null
```

SQL query to generate the prediction input

```
1 select price from houses where quality is null
```

Save Model

TODO

- Automatic syntax check of the python code
- Automatic check of provided model (signature of train/predict functions)
- inputs should be auto-resizeable
- Add model templates
- Add delete/duplicate button

Model templates

Backend

Web server

python3 with flask, based on [microblog](#). For code textboxes we use [CodeMirror](#).

Models

Saved in a sqlite DB. They contain all the text fields from the webpage, along with the trained model state (pickle) and the last result of a train / predict run.

Jobs / Tasks

A task is a background task that we give a worker to do. The workers are implemented using Redis.

train_model

Trains a specific model. The job will execute a sql query, train the model and save (pickle) the trained model in a DB.

run_model

Runs a specific model over some input. The job will execute a sql query, load the model, run it on the sql data, and finally save it to another sql table: <classifier_name>_prediction. We may want to let the user modify the table name.

Saving model state

We need a method of saving/loading the model state. We use python's pickle.dumps() and save it to the models DB. A small catch is that class attributes are **NOT** pickled. Read more [\[here\]](#).

The following, for example, will not work:

```
class Model:
    m = LinearRegression()
    def train(self, X, y):
        self.m.fit(X, y)
    def predict(self, X):
        return self.m.predict(X)
```

This is okay, though:

```
class Model:
    def __main__(self):
        self.m = LinearRegression()
```

```
def train(self, X, y):
    self.m.fit(X, y)
def predict(self, X):
    return self.m.predict(X)
```

Another option is [sklearn's dump\(\)](#), but it only works on sklearn models.

As the model class is dynamically generated, it can not be easily pickled. Using a small [hack](#), we can pickle and unpickle it:

```
names = {}
exec(model_str, names)
global Model
Model = names['Model']
Model.__module__ = __name__
obj = Model()
obj.train(X, y)
state = pickle.dumps(obj)
```

And to unpickle:

```
names = {}
exec(model.ml_model, names)
global Model
Model = names['Model']
Model.__module__ = __name__
obj = pickle.loads(model.state_pk1)
y = obj.predict(X)
```

Predefined Models

Linear Regression

```
from sklearn.linear_model import LinearRegression

class Model:
    def __init__(self):
        self.m = LinearRegression()
    def train(self, X, y):
        self.m.fit(X, y)
    def predict(self, X):
        return self.m.predict(X)
```

Usage

1. Get [code](#) and install dependencies from requirements.txt
2. Start redis server:
\$ redis-server
3. Start a worker (within environment):
\$ rq worker
4. Start webserver:
\$ FLASK_APP=soft_query_app.py DEBUG_MODE=1 flask run
5. Profit

Other

Redis

[Installing redis on OSX](#)

1. Install redis:
\$ brew install redis
\$ brew services start redis
2. Install python redis:
\$ pip3 install rq
3. Start redis server:
\$ redis-server
4. Start a worker (within environment):
\$ rq worker
5. Submit job through python:

```
from redis import Redis
from rq import Queue
import time

q = Queue(connection=Redis())

from rq_test_func import delayed_return
job = q.enqueue(delayed_return, 1, 'http://nvie.com', timeout=1)
print(job.result)
time.sleep(2)
print(job.result)
```

6. Run the above code:
\$ python ./test.py

7. View worker status:

```
$ rq info
```

Possible use case

```
-- Training Input
select user.id, user.embedding, user.cocacola_label -- we need the name of the
user.id column, to create the new table later but not really all the IDs.
where cocacola_label <> ''
from users

-- Inference Input
-- Export stuff to table "<classifier_name>_prediction" with a schema sort of like
(user_id, prediction_result)

select user.id, user.embedding
where cocacola_label == ''
```